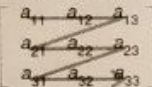
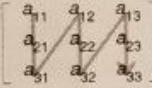


Homework #3-1

Problem 1

Assume memory is byte addressable and words are 64 bits, unless specified otherwise.

The following reference from Wikipedia may be useful in answering this question:

<p>Row-major order</p>  <p>Column-major order</p>  <p>Illustration of row- and column-major order</p>	<p>Row-major order, e.g., C (0-indexed)</p> <table><thead><tr><th>Address</th><th>Access</th><th>Value</th></tr></thead><tbody><tr><td>0</td><td>A[0][0]</td><td>a₁₁</td></tr><tr><td>1</td><td>A[0][1]</td><td>a₁₂</td></tr><tr><td>2</td><td>A[0][2]</td><td>a₁₃</td></tr><tr><td>3</td><td>A[1][0]</td><td>a₂₁</td></tr><tr><td>4</td><td>A[1][1]</td><td>a₂₂</td></tr><tr><td>5</td><td>A[1][2]</td><td>a₂₃</td></tr></tbody></table>	Address	Access	Value	0	A[0][0]	a ₁₁	1	A[0][1]	a ₁₂	2	A[0][2]	a ₁₃	3	A[1][0]	a ₂₁	4	A[1][1]	a ₂₂	5	A[1][2]	a ₂₃	<p>Column-major order, e.g., Fortran (1-indexed)</p> <table><thead><tr><th>Address</th><th>Access</th><th>Value</th></tr></thead><tbody><tr><td>1</td><td>A(1,1)</td><td>a₁₁</td></tr><tr><td>2</td><td>A(2,1)</td><td>a₂₁</td></tr><tr><td>3</td><td>A(1,2)</td><td>a₁₂</td></tr><tr><td>4</td><td>A(2,2)</td><td>a₂₂</td></tr><tr><td>5</td><td>A(1,3)</td><td>a₁₃</td></tr><tr><td>6</td><td>A(2,3)</td><td>a₂₃</td></tr></tbody></table>	Address	Access	Value	1	A(1,1)	a ₁₁	2	A(2,1)	a ₂₁	3	A(1,2)	a ₁₂	4	A(2,2)	a ₂₂	5	A(1,3)	a ₁₃	6	A(2,3)	a ₂₃
Address	Access	Value																																										
0	A[0][0]	a ₁₁																																										
1	A[0][1]	a ₁₂																																										
2	A[0][2]	a ₁₃																																										
3	A[1][0]	a ₂₁																																										
4	A[1][1]	a ₂₂																																										
5	A[1][2]	a ₂₃																																										
Address	Access	Value																																										
1	A(1,1)	a ₁₁																																										
2	A(2,1)	a ₂₁																																										
3	A(1,2)	a ₁₂																																										
4	A(2,2)	a ₂₂																																										
5	A(1,3)	a ₁₃																																										
6	A(2,3)	a ₂₃																																										
		Matlab is also column-major ordered.																																										
Ref: "Row- and column-major order," Wikipedia. 25-Mar-2018.																																												

- 5.1 In this exercise we look at memory locality properties of matrix computation. The following code is written in C, where elements within the same row are stored contiguously. Assume each word is a 64-bit integer.

```
for (I=0; I<8; I++)
  for (J=0; J<8000; J++)
    A[I][J]=B[I][0]+A[J][I];
```

of bits in 16 byte cache line = $16 \times 8 = 128$
 # of 64 bit integer in 16 byte cache line = $128 / 64 = 2$

- 5.1.1 How many 64-bit integers can be stored in a 16-byte cache block? 2.
- 5.1.2 Which variable references exhibit temporal locality? I and J are constantly accessed.
- 5.1.3 Which variable references exhibit spatial locality? $A[I][J]$. Multiple of them are loaded into cache on a single miss. As I is incremented, nearby values of array are accessed.
- Because of these variables are likely to stay in cache and hence exhibit temporal locality. $B[I][0]$ exhibit temporal locality since I is accessed over and over in a row. $\% I, J, B[I][0]$

Locality is affected by both the reference order and data layout. The same computation can also be written below in Matlab, which differs from C in that it stores matrix elements within the same column contiguously in memory.

```
for I=1:8
    for J=1:8000
        A(I,J)=B(I,0)+A(J,I);
    end
end
```

• I and J are accessed every iteration of the loop. It will stay in cache because they are exhibiting temporal locality.

5.1.4 Which variable references exhibit temporal locality? $B[I][0], I, J$

5.1.5 Which variable references exhibit spatial locality? $A[J, I], B[I][0]$

5.1.6 How many 16-byte cache blocks are needed to store all 64-bit matrix elements being referenced using Matlab's matrix storage?

1st. 2nd.
 $A[I, J]$ for values of I ranging between 0 and 9999 and values of J ranging between 0 and 7.

There are $8000 \times 8 = 64000$

$64000 \times 2 = 128000$

! There is a rectangle of $8 \times 8 = 64$ element that is common between both groups of elements and should be subtracted.

$128000 - 64 = 127936$

$127936 + 8 = 127944$ $127936 + 8 = 127944$

*there are 8 elements being accessed from any B.
 Total number of elements*

Since each cache line can store 4 elements,

we would need $127944 / 4 = 31986$ cache lines

to store the entire matrix.

Problem 2

5.5 For a direct-mapped cache design with a 64-bit address, the following bits of the address are used to access the cache.

Tag	Index	Offset
63-10	9-5	4-0

→ how many blocks are present in directly mapped cache.

5.5.1 What is the cache block size (in words)? 0-4 offset. $2^5 = 32$. Each cache block has four 8-byte words.

5.5.2 How many blocks does the cache have? 9-5 Index. $2^5 = 32$

5.5.3 What is the ratio between total bits required for such a cache implementation over the data storage bits? 54 bits for each block tag field. $\frac{2^5 + 216 + 1240}{2^{10} + 1024} = 1.21$
 With 32 blocks, $54 \times 32 = 1728$ bits for tag field.
 Total bits for the cache = $2^{10} + 1728 = 1024 + 1728 = 2752$ bits.

Beginning from power on, the following byte-addressed cache references are recorded.

Address											
Hex	00	04	10	84	E8	A0	400	1E	8C	C1C	B4
Dec	0	4	16	132	232	160	1024	30	140	3100	180
	M	H	H	M	M	M	M	M	H	M	H

Annotations:

- Mandatory Miss first time referenced
- Index and tag are same as 0x00
- Same index, different tag as 0x04
- Same index as 4th frame
- same index, different tag as 0x1E
- Index and tag are same as 0x04
- Mandatory Miss
- Same index, diff tag as 0x400
- Same index as 0x A0
- the frame needs replacement

5.5.4 For each reference, list:

1. Its tag and offset;
2. Whether it is a hit or a miss;
3. Which bytes were replaced (if any).

5.5.5 What is the hit ratio? $\frac{4}{12} = 0.33 \times 100 = 33\%$

00	0000 0000 0000 0000
04	0000 0000 0000 0100
10	0000 0000 0000 0000
84	0000 0000 1000 0100
E8	0000 0000 1110 1000
A0	0000 0000 1010 0000
400	0000 0100 0000 0000
1E	0000 0000 0001 1110
8C	0000 0000 1000 1100
C1C	0000 1100 0001 1100
B4	0000 0000 1011 0100
884	0000 1000 1000 0100

Problem 3

5.6 Recall that we have two write policies and two write allocation policies, and their combinations can be implemented either in L1 or L2 cache. Assume the following choices for L1 and L2 caches:

L1	L2
Write through, non-write allocate	Write back, write allocate

5.6.1 Buffers are employed between different levels of memory hierarchy to reduce access latency. For this given configuration, list the possible buffers needed between L1 and L2 caches, as well as L2 cache and memory. *L1 \Rightarrow write*

5.6.2 Describe the procedure of handling an L1 write-miss, considering the components involved and the possibility of replacing a dirty block.

5.6.3 For a multilevel exclusive cache configuration (a block can only reside in one of the L1 and L2 caches), describe the procedures of handling an L1 write-miss and an L1 read-miss, considering the components involved and the possibility of replacing a dirty block.

5.6.1. Between L1 and L2 caches, one write buffer is required. When the miss occurs, we directly update the portion of the block into the buffer, which will be waiting to be written into L2 cache, while the process doesn't need to stall if the buffer is not full.

Between L2 cache and the memory, we require write and store buffers. When we have a cache miss, we must first write the block back to memory. If the data in the cache is modified. In this situation, a write buffer is required to hold that data, such that the processor can continue the execution while that data is waiting to be written to the memory. In the meanwhile, a store buffer is used, such that the processor places the new data in the store buffer. Then ~~now~~ when a cache hit occurs, this new data is written from the store buffer into the cache.

5.6.2. First we check whether the block is dirty. If it is, then we write the dirty block to memory. Next, we retrieve the target block from memory (overwriting the block that is in our way). Finally we write to our L2 block.